

Pollution SARIMAX TimeSeries Forecasting

May 28, 2019

```
In [1]: !pip install --user xlrd
```

Requirement already satisfied: xlrd in /home/omar/.local/lib/python3.5/site-packages (1.2.0)
You are using pip version 19.0.3, however version 19.1.1 is available. You should consider upgrading

```
In [2]: !ls
```

```
ARIMA TimeSeries Forecasting-Copy2.ipynb
ARIMA TimeSeries Forecasting-Copy3.ipynb
ARIMA TimeSeries Forecasting.ipynb
ARIMA TimeSeries Forecasting.pdf
cleaned_data
eda and time series prediction.ipynb
Historical Pollen Index
Historical Pollution Index
Historical Symptom Logs
__MACOSX
method_1_ARIMA_timeseries_prediction_sample.ipynb
pollen.csv
pollen_history.csv
SARIMAX Pollen TimeSeries Forecasting.ipynb
SARIMAX Pollution TimeSeries Forecasting.ipynb
symptom_cause_2019051616-03_001.csv
weatherStored_2019051615-59_001.csv
```

```
In [3]: !ls cleaned_data/
```

```
Pollen_full_data_Eve.xlsx      Symptom Cause Eve.xlsx
Pollution full data Eve.xlsx  Symptom log full data Eve.xlsx
~$Symptom Cause Eve.xlsx
```

```
In [4]: import pandas as pd
        from datetime import datetime

        import plotly
```

```

plotly.offline.init_notebook_mode()

import plotly.plotly as py
import plotly.graph_objs as go

from statsmodels.tsa.stattools import adfuller
from numpy import log

import itertools
import numpy as np
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")

plt.style.use('fivethirtyeight')
import pandas as pd
import statsmodels.api as sm
import matplotlib
matplotlib.rcParams['axes.labelsize'] = 14
matplotlib.rcParams['xtick.labelsize'] = 12
matplotlib.rcParams['ytick.labelsize'] = 12
matplotlib.rcParams['text.color'] = 'k'

```

```
In [5]: data = pd.read_excel('cleaned_data/Pollution full data Eve.xlsx')
```

```
In [6]: list(data)
```

```

Out[6]: ['city',
         'date_time',
         'date_time.1',
         'pollutantGlobalIndex',
         'so2Index',
         'no2Index',
         'o3Index',
         'coIndex',
         'pm25Index',
         'pm10Index']

```

```
In [7]: data.head()
```

```

Out[7]:
   city  date_time  date_time.1  pollutantGlobalIndex \
0  Austin 2017-09-06 14:00:00 2017-09-06 14:00:00      3
1  Austin 2017-09-06 15:00:00 2017-09-06 15:00:00      3
2  Austin 2017-09-06 16:00:00 2017-09-06 16:00:00      3
3  Austin 2017-09-06 17:00:00 2017-09-06 17:00:00      3
4  Austin 2017-09-06 18:00:00 2017-09-06 18:00:00      3

   so2Index  no2Index  o3Index  coIndex  pm25Index  pm10Index

```

0	1	1	3	NaN	3	NaN
1	1	1	3	NaN	3	NaN
2	1	1	3	NaN	3	NaN
3	1	1	3	NaN	2	NaN
4	1	1	3	NaN	2	NaN

```
In [8]: data = data[['date_time', 'pollutantGlobalIndex']]
```

```
In [9]: data.head()
```

```
Out[9]:
```

	date_time	pollutantGlobalIndex
0	2017-09-06 14:00:00	3
1	2017-09-06 15:00:00	3
2	2017-09-06 16:00:00	3
3	2017-09-06 17:00:00	3
4	2017-09-06 18:00:00	3

```
In [11]: data.dropna()
```

```
Out[11]:
```

	date_time	pollutantGlobalIndex
0	2017-09-06 14:00:00	3
1	2017-09-06 15:00:00	3
2	2017-09-06 16:00:00	3
3	2017-09-06 17:00:00	3
4	2017-09-06 18:00:00	3
5	2017-09-06 19:00:00	3
6	2017-09-06 20:00:00	3
7	2017-09-06 21:00:00	3
8	2017-09-06 22:00:00	3
9	2017-09-06 23:00:00	3
10	2017-09-07 00:00:00	3
11	2017-09-07 01:00:00	3
12	2017-09-07 02:00:00	4
13	2017-09-07 03:00:00	4
14	2017-09-07 04:00:00	4
15	2017-09-07 05:00:00	4
16	2017-09-07 06:00:00	4
17	2017-09-07 07:00:00	4
18	2017-09-07 08:00:00	4
19	2017-09-07 09:00:00	4
20	2017-09-07 10:00:00	3
21	2017-09-07 11:00:00	3
22	2017-09-07 12:00:00	3
23	2017-09-07 13:00:00	3
24	2017-09-07 14:00:00	3
25	2017-09-07 15:00:00	3
26	2017-09-07 16:00:00	3
27	2017-09-07 17:00:00	2
28	2017-09-07 18:00:00	2

29	2017-09-07 19:00:00	2
...
26244	2019-05-20 07:00:00	6
26245	2019-05-20 08:00:00	6
26246	2019-05-20 09:00:00	6
26247	2019-05-20 10:00:00	6
26248	2019-05-20 11:00:00	6
26249	2019-05-20 12:00:00	7
26250	2019-05-20 13:00:00	7
26251	2019-05-20 14:00:00	7
26252	2019-05-20 15:00:00	7
26253	2019-05-20 16:00:00	8
26254	2019-05-20 17:00:00	9
26255	2019-05-20 18:00:00	8
26256	2019-05-20 19:00:00	8
26257	2019-05-20 20:00:00	7
26258	2019-05-20 21:00:00	7
26259	2019-05-20 22:00:00	7
26260	2019-05-20 23:00:00	7
26261	2019-05-21 00:00:00	6
26262	2019-05-21 01:00:00	4
26263	2019-05-21 02:00:00	4
26264	2019-05-21 03:00:00	4
26265	2019-05-21 04:00:00	4
26266	2019-05-21 05:00:00	6
26267	2019-05-21 06:00:00	4
26268	2019-05-21 07:00:00	7
26269	2019-05-21 08:00:00	6
26270	2019-05-21 09:00:00	7
26271	2019-05-21 10:00:00	7
26272	2019-05-21 11:00:00	7
26273	2019-05-21 12:00:00	7

[26274 rows x 2 columns]

```
print(data.index.min() + " " + data.index.max())
a = data['date_time'].astype(str).str.split(" ", n = 1, expand = True)
data['date'] = a[0]
```

In [12]: data.head()

```
Out[12]:
```

	date_time	pollutantGlobalIndex
0	2017-09-06 14:00:00	3
1	2017-09-06 15:00:00	3
2	2017-09-06 16:00:00	3
3	2017-09-06 17:00:00	3
4	2017-09-06 18:00:00	3

In [14]: data = data.groupby('date_time')['pollutantGlobalIndex'].sum().reset_index()

```
In [15]: data.head(20)
```

```
Out[15]:
```

	date_time	pollutantGlobalIndex
0	2017-09-06 14:00:00	3
1	2017-09-06 15:00:00	3
2	2017-09-06 16:00:00	3
3	2017-09-06 17:00:00	3
4	2017-09-06 18:00:00	3
5	2017-09-06 19:00:00	3
6	2017-09-06 20:00:00	3
7	2017-09-06 21:00:00	3
8	2017-09-06 22:00:00	3
9	2017-09-06 23:00:00	3
10	2017-09-07 00:00:00	3
11	2017-09-07 01:00:00	3
12	2017-09-07 02:00:00	4
13	2017-09-07 03:00:00	4
14	2017-09-07 04:00:00	4
15	2017-09-07 05:00:00	4
16	2017-09-07 06:00:00	4
17	2017-09-07 07:00:00	4
18	2017-09-07 08:00:00	4
19	2017-09-07 09:00:00	4

```
data['2017-09-07'].max()
```

```
In [18]: type(data)
```


```
Out[18]: pandas.core.frame.DataFrame
```

```
In [19]: data.set_index('date_time', inplace=True)
```

```
In [20]: data.index = pd.to_datetime(data.index)
```

```
In [21]: data.plot()
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1a9ff1a2b0>
```



We will use a “grid search” to iteratively explore different combinations of parameters. For each combination of parameters, we fit a new seasonal ARIMA model with the SARIMAX() function from the statsmodels module and assess its overall quality. Once we have explored the entire landscape of parameters, our optimal set of parameters will be the one that yields the best performance for our criteria of interest. Let’s begin by generating the various combination of parameters that we wish to assess:

```
In [22]: # Define the p, d and q parameters to take any value between 0 and 2
p = d = q = range(0, 2)

# Generate all different combinations of p, q and q triplets
pdq = list(itertools.product(p, d, q))

# Generate all different combinations of seasonal p, q and q triplets
seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]

print('Examples of parameter combinations for Seasonal ARIMA...')
print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[1]))
print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[2]))
print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[3]))
print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[4]))
```

Examples of parameter combinations for Seasonal ARIMA...

SARIMAX: (0, 0, 1) x (0, 0, 1, 12)

SARIMAX: (0, 0, 1) x (0, 1, 0, 12)

```
SARIMAX: (0, 1, 0) x (0, 1, 1, 12)
SARIMAX: (0, 1, 0) x (1, 0, 0, 12)
```

```
In [23]: for param in pdq:
          for param_seasonal in seasonal_pdq:
              try:
                  mod = sm.tsa.statespace.SARIMAX(y,
                                                    order=param,
                                                    seasonal_order=param_seasonal,
                                                    enforce_stationarity=False,
                                                    enforce_invertibility=False)

                  results = mod.fit()

                  print('ARIMA{0}x{1}12 - AIC:{0}'.format(param, param_seasonal, results.aic))
              except:
                  continue
```

```
In [24]: y = data
```

The code below iterates through combinations of parameters and uses the SARIMAX function from statsmodels to fit the corresponding Seasonal ARIMA model. Here, the order argument specifies the (p, d, q) parameters, while the seasonal_order argument specifies the (P, D, Q, S) seasonal component of the Seasonal ARIMA model. After fitting each SARIMAX() model, the code prints out its respective AIC score.

The lesser the AIC score, the better model architecture it is.

Akaike's Information Criterion (AIC): Formally, AIC is defined as $2\log L + 2k$ where L is the maximized log likelihood and k is the number of parameters in the model. For the normal regression problem, AIC is an estimate of the Kullback-Leibler discrepancy between a true model and a candidate model.

```
In [25]: for param in pdq:
          for param_seasonal in seasonal_pdq:

              mod = sm.tsa.statespace.SARIMAX(y,
                                                order=param,
                                                seasonal_order=param_seasonal,
                                                enforce_stationarity=False,
                                                enforce_invertibility=False)

              results = mod.fit()

              print('ARIMA{0}x{1}12 - AIC:{0}'.format(param, param_seasonal, results.aic))

ARIMA(0, 0, 0)x(0, 0, 0, 12)12 - AIC:106687.52803499259
ARIMA(0, 0, 0)x(0, 0, 1, 12)12 - AIC:95806.76451832935
ARIMA(0, 0, 0)x(0, 1, 0, 12)12 - AIC:77719.46114448212
ARIMA(0, 0, 0)x(0, 1, 1, 12)12 - AIC:72595.83401805462
```

ARIMA(0, 0, 0)x(1, 0, 0, 12)12 - AIC:77227.17792550023
 ARIMA(0, 0, 0)x(1, 0, 1, 12)12 - AIC:72622.20617441263
 ARIMA(0, 0, 0)x(1, 1, 0, 12)12 - AIC:73828.76874015582
 ARIMA(0, 0, 0)x(1, 1, 1, 12)12 - AIC:72436.46745807579
 ARIMA(0, 0, 1)x(0, 0, 0, 12)12 - AIC:90587.40701959425
 ARIMA(0, 0, 1)x(0, 0, 1, 12)12 - AIC:82460.52020150796
 ARIMA(0, 0, 1)x(0, 1, 0, 12)12 - AIC:70112.29216662335
 ARIMA(0, 0, 1)x(0, 1, 1, 12)12 - AIC:64026.149056888775
 ARIMA(0, 0, 1)x(1, 0, 0, 12)12 - AIC:69393.15984711159
 ARIMA(0, 0, 1)x(1, 0, 1, 12)12 - AIC:64053.58289285139
 ARIMA(0, 0, 1)x(1, 1, 0, 12)12 - AIC:66009.8563356791
 ARIMA(0, 0, 1)x(1, 1, 1, 12)12 - AIC:63922.36949658996
 ARIMA(0, 1, 0)x(0, 0, 0, 12)12 - AIC:56126.59866557239
 ARIMA(0, 1, 0)x(0, 0, 1, 12)12 - AIC:56091.320396223324
 ARIMA(0, 1, 0)x(0, 1, 0, 12)12 - AIC:65918.84910609439
 ARIMA(0, 1, 0)x(0, 1, 1, 12)12 - AIC:56117.767094150986
 ARIMA(0, 1, 0)x(1, 0, 0, 12)12 - AIC:56093.867617736585
 ARIMA(0, 1, 0)x(1, 0, 1, 12)12 - AIC:56067.39256360466
 ARIMA(0, 1, 0)x(1, 1, 0, 12)12 - AIC:61173.373579326944
 ARIMA(0, 1, 0)x(1, 1, 1, 12)12 - AIC:56119.73748484236
 ARIMA(0, 1, 1)x(0, 0, 0, 12)12 - AIC:55599.74150340479
 ARIMA(0, 1, 1)x(0, 0, 1, 12)12 - AIC:55565.48645786028
 ARIMA(0, 1, 1)x(0, 1, 0, 12)12 - AIC:65416.96440023401
 ARIMA(0, 1, 1)x(0, 1, 1, 12)12 - AIC:55585.682663974265
 ARIMA(0, 1, 1)x(1, 0, 0, 12)12 - AIC:55570.91907491916
 ARIMA(0, 1, 1)x(1, 0, 1, 12)12 - AIC:55519.595768934116
 ARIMA(0, 1, 1)x(1, 1, 0, 12)12 - AIC:60535.490496210725
 ARIMA(0, 1, 1)x(1, 1, 1, 12)12 - AIC:55587.66450202813
 ARIMA(1, 0, 0)x(0, 0, 0, 12)12 - AIC:56020.82034273728
 ARIMA(1, 0, 0)x(0, 0, 1, 12)12 - AIC:55982.410360225986
 ARIMA(1, 0, 0)x(0, 1, 0, 12)12 - AIC:64227.05133302646
 ARIMA(1, 0, 0)x(0, 1, 1, 12)12 - AIC:55348.879253294566
 ARIMA(1, 0, 0)x(1, 0, 0, 12)12 - AIC:55981.549733660904
 ARIMA(1, 0, 0)x(1, 0, 1, 12)12 - AIC:55373.61039903519
 ARIMA(1, 0, 0)x(1, 1, 0, 12)12 - AIC:59587.17551528601
 ARIMA(1, 0, 0)x(1, 1, 1, 12)12 - AIC:55350.87204135413
 ARIMA(1, 0, 1)x(0, 0, 0, 12)12 - AIC:55535.51708667213
 ARIMA(1, 0, 1)x(0, 0, 1, 12)12 - AIC:55499.60586893518
 ARIMA(1, 0, 1)x(0, 1, 0, 12)12 - AIC:64150.15088835024
 ARIMA(1, 0, 1)x(0, 1, 1, 12)12 - AIC:55102.36067543709
 ARIMA(1, 0, 1)x(1, 0, 0, 12)12 - AIC:55501.85792065177
 ARIMA(1, 0, 1)x(1, 0, 1, 12)12 - AIC:55124.72102552027
 ARIMA(1, 0, 1)x(1, 1, 0, 12)12 - AIC:59444.75120010918
 ARIMA(1, 0, 1)x(1, 1, 1, 12)12 - AIC:55104.36054931962
 ARIMA(1, 1, 0)x(0, 0, 0, 12)12 - AIC:55641.03146139863
 ARIMA(1, 1, 0)x(0, 0, 1, 12)12 - AIC:55606.35577679375
 ARIMA(1, 1, 0)x(0, 1, 0, 12)12 - AIC:65442.34636216519
 ARIMA(1, 1, 0)x(0, 1, 1, 12)12 - AIC:55628.64166917707


```

ARIMA(1, 1, 0)x(1, 0, 0, 12)12 - AIC:55606.02325671892
ARIMA(1, 1, 0)x(1, 0, 1, 12)12 - AIC:55565.64356661648
ARIMA(1, 1, 0)x(1, 1, 0, 12)12 - AIC:60600.26351396194
ARIMA(1, 1, 0)x(1, 1, 1, 12)12 - AIC:55630.6163411032
ARIMA(1, 1, 1)x(0, 0, 0, 12)12 - AIC:54951.932862377056
ARIMA(1, 1, 1)x(0, 0, 1, 12)12 - AIC:54913.5231834527
ARIMA(1, 1, 1)x(0, 1, 0, 12)12 - AIC:64231.060852482886
ARIMA(1, 1, 1)x(0, 1, 1, 12)12 - AIC:54945.63165784913
ARIMA(1, 1, 1)x(1, 0, 0, 12)12 - AIC:54915.94299261454
ARIMA(1, 1, 1)x(1, 0, 1, 12)12 - AIC:54889.8878570757
ARIMA(1, 1, 1)x(1, 1, 0, 12)12 - AIC:59595.613022084566
ARIMA(1, 1, 1)x(1, 1, 1, 12)12 - AIC:54947.548306623394

```

```

In [26]: mod = sm.tsa.statespace.SARIMAX(y,
                                         order=(1, 1, 1),
                                         seasonal_order=(1, 1, 1, 12),
                                         enforce_stationarity=False,
                                         enforce_invertibility=False)

results = mod.fit()

print(results.summary().tables[1])

```

```

=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.7840	0.004	183.397	0.000	0.776	0.792
ma.L1	-1.0366	0.002	-452.991	0.000	-1.041	-1.032
ar.S.L12	0.0003	0.006	0.050	0.960	-0.011	0.012
ma.S.L12	-0.9980	0.001	-738.487	0.000	-1.001	-0.995
sigma2	2.4233	0.015	157.330	0.000	2.393	2.453

```

=====

```

```

In [27]: results.plot_diagnostics(figsize=(15, 12))
plt.show()

```